

Deploy a web application in a cloud-based Kubernetes solution, ensuring proper logging and monitoring are in place.

Application Functionality

The web application shall have a simple static page of content.

Through Terraform we will deploy - VPC, EKS Cluster, Node Group, IAM Role associated with eks and nodegroup.

Create a folder and inside the folder create .tf file for -

- provider.tf - We need to specify terraform the resources it need to interact and manage, here i have provided aws provider as we are deploying resources in AWS Environment.
- vpc.tf - Need to create a vpc, where we can deploy our eks cluster and worker nodes. I have created 2 public and 2 private subnets.
- eks-cluster.tf - Inside the file we need to specify terraform about the creation of resources - EKS cluster, nodegroups, IAM policy, IAM role and attachment with resources created for permissions.
- variables.tf - The variables used within the above templates, here we can provide values to them.

We can also create terraform workspace for multi-environment configuration, helping in maintaining particular state file. We can use .tfvars file to provide the value to the variables used.

To create webserver, we are using nginx which we will specify in Dockerfile. Make sure to create the image as per the architecture of worker nodes otherwise it will create an error at the time of deployment of pods. For content to show in webpage, create an index.html file and provide valid content. Now create an image through docker build image command and push the image to docker hub. I have provided the image publicly available. Use the image inside the spec of container in deployment. It will build the pods with that container image.

The screenshot shows the Docker Hub interface for the repository 'ashish13997/web-app'. The page includes a sidebar with navigation options like Repositories, Settings, and Billing. The main content area displays the repository details, including the name 'ashish13997/web-app', the last push time 'about 2 hours ago', and the repository size '35.9 MB'. There are tabs for General, Tags, Image Management, Collaborators, Webhooks, and Settings. The 'Tags' tab is selected, showing a list of tags. The first tag is 'v1', which was last pushed about 2 hours ago by 'ashish13997'. Below the tag list, there is a table with columns for Digest, OS/ARCH, Last pull, and Compressed size. The table shows the digest 'a175d7cfba64' for the 'linux/amd64' architecture, with the last pull occurring 'less than 1 day' ago and a compressed size of '16.07 MB'. On the right side, there are Docker commands for pushing and pulling the image, and a 'Public view' button.

Digest	OS/ARCH	Last pull	Compressed size
a175d7cfba64	linux/amd64	less than 1 day	16.07 MB

In Kubernetes manifest, we will create yaml file for -

- namespace.yaml - To isolate and organise the resources with others. Basically it act as environment where we can create multi namespaces to differentiate prod, dev and provide limits with resource quota.
- deployment.yaml - This file help us with managing and deploying our application. Here i have specified replicas = 2, so we will have two pods created with specified image within spec of container's template. In labels we have to specify namespace, as per the value it will deploy the template under that namespace.
- service.yaml - This file help us in exposing our application to network traffic. Here i have taken Load Balancer, it is best practise to deploy worker nodes in private subnets. Therefore, to access our webserver, load balancer is used to provide internet access. By default, in service if we opt for load balancer it will take classic load balancer.
- cloudwatch-fluentd.yaml - To monitor our eks cluster and nodes cpu utilisation, memory etc. we are using AWS CloudWatch Container-insights, through the template we are creating namespace for cloudwatch, security account, cluster role, cluster role binding, daemon set for cloudwatch agent and fluentd to collect metrics and send them to AWS container insights, config-map for both the daemon sets.

Steps -

- After the creation of above templates, give command - terraform init
- To initialise the mandate plugins as specified under provider.tf file.
- To check the output of our terraform file, give command: terraform plan, it will give us an output regarding resource creation.
- After checking give command: terraform plan and yes or —auto-approve
- After the creation of VPC, EKS Cluster, nodegroups and IAM Role.
- Verify Namespace - kubectl get ns
- Verify EKS Cluster - aws eks list-clusters
- To describe the cluster - aws eks describe-cluster --name test-eks-cluster
- Verify worker nodes are up and running - kubectl get nodes

The screenshot shows a VS Code editor with several Terraform files open: 02-deployment.yaml, Dockerfile, index.html, 03-service.yaml, eks-cluster.tf, 01-namespace.yaml, and 04-cwagent-fluentd-qu. The main editor window displays the content of eks-cluster.tf, which defines an IAM role and its policy for the EKS cluster. The terminal window at the bottom shows the output of a terraform plan command, indicating that a VPC will be created with the name 'test-eks-cluster-vpc' and tagged with 'dev', 'test-eks-cluster-vpc', and 'test-eks-cluster'. The plan also shows the creation of an IAM role and its policy attachment. The terminal output includes the following details:

```
Plan: 24 to add, 0 to change, 0 to destroy.

Changes to Outputs:
+ cluster_endpoint = (known after apply)
+ cluster_name     = "test-eks-cluster"

Do you want to perform these actions?
Terraform will perform the actions described above.
Only 'yes' will be accepted to approve.

Enter a value: yes

aws_iam_role.eks_cluster_role: Creating...
```

```

aws_eks_node_group.main: Still creating... [1m10s elapsed]
aws_eks_node_group.main: Still creating... [1m20s elapsed]
aws_eks_node_group.main: Still creating... [1m30s elapsed]
aws_eks_node_group.main: Still creating... [1m40s elapsed]
aws_eks_node_group.main: Creation complete after 1m49s [id=test-eks-cluster:test-eks-cluster-node-group]

Apply complete! Resources: 24 added, 0 changed, 0 destroyed.

● Outputs:

cluster_endpoint = "https://FE5E40534081E2727F92B0C8B9FF96D7.gr7.ap-south-1.eks.amazonaws.com"
cluster_name = "test-eks-cluster"
tanish@Tanishs-Laptop ATC-Assignment % aws eks update-kubeconfig --name test-eks-cluster --region ap-south-1
Updated context arn:aws:eks:ap-south-1:767398024481:cluster/test-eks-cluster in /Users/tanish/.kube/config
● tanish@Tanishs-Laptop ATC-Assignment % kubectl get nodes
NAME                                     STATUS    ROLES    AGE     VERSION
ip-10-0-1-70.ap-south-1.compute.internal Ready    <none>    5m44s   v1.28.15-eks-aeac579
ip-10-0-2-67.ap-south-1.compute.internal Ready    <none>    5m46s   v1.28.15-eks-aeac579
tanish@Tanishs-Laptop ATC-Assignment %

```

```

● tanish@Tanishs-Laptop ATC-Assignment % aws eks list-clusters
{
  "clusters": [
    "test-eks-cluster"
  ]
}
● tanish@Tanishs-Laptop ATC-Assignment % kubectl get nodes
NAME                                     STATUS    ROLES    AGE     VERSION
ip-10-0-1-70.ap-south-1.compute.internal Ready    <none>    7m42s   v1.28.15-eks-aeac579
ip-10-0-2-67.ap-south-1.compute.internal Ready    <none>    7m44s   v1.28.15-eks-aeac579
○ tanish@Tanishs-Laptop ATC-Assignment %

```

- Go the path where our kubernetes manifest folder is present, inside manifest folder add the namespace, deployment, service and cloud-watch yaml files.
- Give command - `cd ..`
- Now give command - `kubectl apply -f <kube-manifest-folder-name>/`
It will automatically create all the yaml files and configuration inside given.

```

tanish@Tanishs-Laptop ATC-Assignment % kubectl apply -f kube-manifests
namespace/test created
deployment.apps/web-app created
service/web-app-service created
namespace/amazon-cloudwatch created
serviceaccount/cloudwatch-agent created
clusterrole.rbac.authorization.k8s.io/cloudwatch-agent-role created
clusterrolebinding.rbac.authorization.k8s.io/cloudwatch-agent-role-binding created
configmap/cwagentconfig created
daemonset.apps/cloudwatch-agent created
configmap/cluster-info created
serviceaccount/fluentd created
clusterrole.rbac.authorization.k8s.io/fluentd-role created
clusterrolebinding.rbac.authorization.k8s.io/fluentd-role-binding created
configmap/fluentd-config created
daemonset.apps/fluentd-cloudwatch created

```

```

tanish@Tanishs-Laptop ATC-Assignment % kubectl get pods -n test
NAME                                READY   STATUS    RESTARTS   AGE
web-app-75b6c67489-2x7h5           1/1     Running   0           72s
web-app-75b6c67489-4hrhv           1/1     Running   0           72s
tanish@Tanishs-Laptop ATC-Assignment % kubectl get pods -n amazon-cloudwatch
NAME                                READY   STATUS    RESTARTS   AGE
cloudwatch-agent-ctd8w             1/1     Running   0           87s
cloudwatch-agent-gkv5l             1/1     Running   0           87s
fluentd-cloudwatch-7bx8j           1/1     Running   0           87s
fluentd-cloudwatch-mhsr6           1/1     Running   0           87s
tanish@Tanishs-Laptop ATC-Assignment % kubectl get svc -n test
NAME                                TYPE           CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
web-app-service                     LoadBalancer  172.20.206.139  ad87341296a82413eb7b5b035cc70047-1731120878.ap-south-1.elb.amazonaws.com  80:31531/TCP     2m

```

- To check pods are running for app server and cloudwatch -
- `kubectl get pods -n amazon-cloudwatch`
- `kubectl get pods -n test`
- To access our webpage - `kubectl get svc -n test`
- Copy the DNS name or from AWS console > EC2 > LoadBalancer > Click on LoadBalancer created and copy the DNS Name.

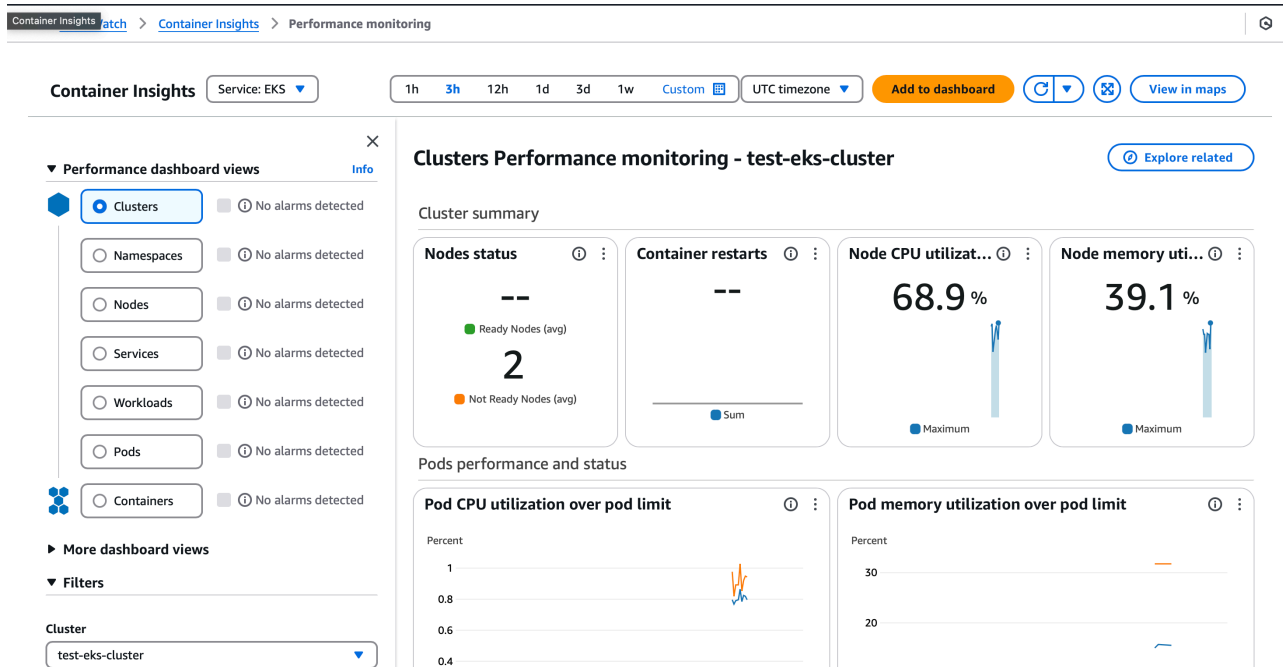
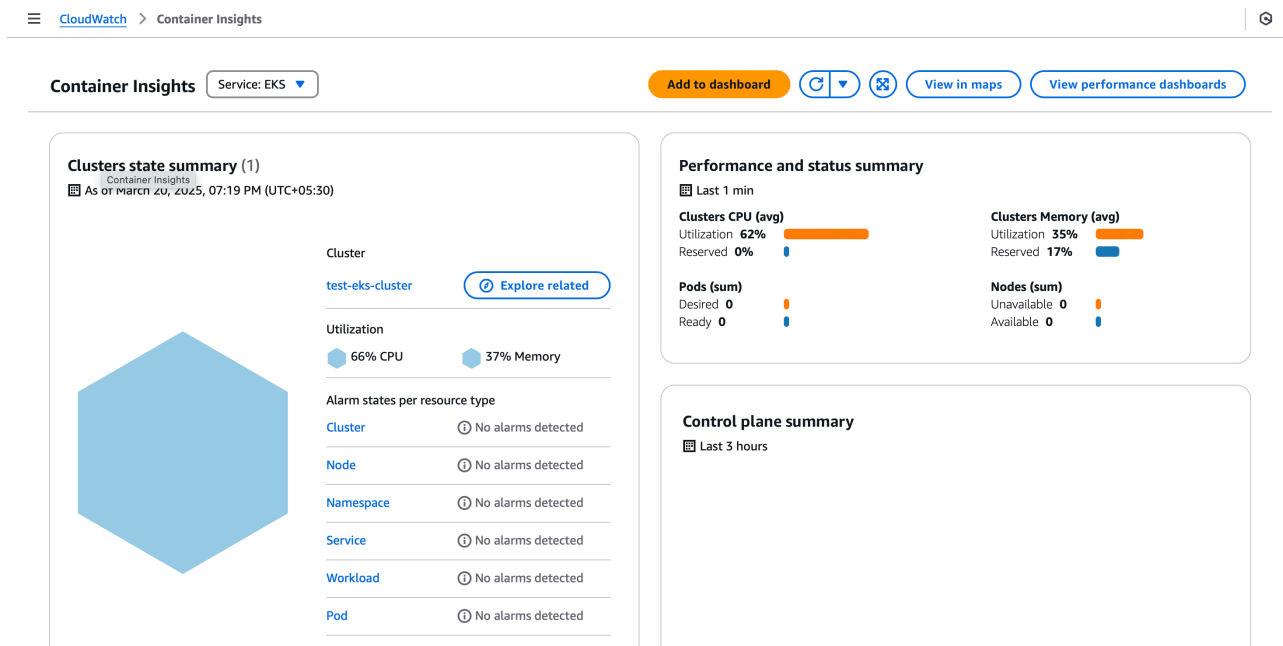
The screenshot shows the AWS Management Console interface. The breadcrumb navigation indicates we are in the EC2 console, specifically under Load balancers, and we have selected a specific load balancer with ID ad87341296a82413eb7b5b035cc70047. The left-hand navigation pane shows various AWS services, with 'Load Balancing' expanded. The main content area displays the configuration for the selected load balancer. Key details include: Load balancer type (Classic), Status (2 of 2 instances in service), Scheme (Internet-facing), Hosted zone (ZP97RAFLXTNZK), VPC (vpc-0410a25573169c3f2), and Availability Zones (subnet-08417ccb75da98afe in ap-south-1b and subnet-05a25abf0eb59eb6c in ap-south-1a). The DNS name is listed as ad87341296a82413eb7b5b035cc70047-1731120878.ap-south-1.elb.amazonaws.com (A Record). Below this, there is a section for 'Distribution of targets by Availability Zone (AZ)' and a series of tabs: Listeners, Network mapping, Security, Health checks, Target instances, Monitoring, and Attribution. The 'Listeners' tab is currently selected.

- Paste the URL over the browser, here we can access the webpage.

The screenshot shows a web browser window. The address bar displays the URL: 14c94a4badc64df95e359da-904543115.ap-south-1.elb.amazonaws.com. The page content is a simple text message: "Hello, welcome to ATC assignment".

Hello, welcome to ATC assignment

- To check the monitoring and utilisation of our cluster and nodes, go to AWS Console > CloudWatch > Container Insights
- Here we can see metrics related to our EKS Cluster.



Based upon our load we can also use Vertical Pod Autoscaler (By default present within EKS) and Horizontal Pod Autoscaler.